

Value-Based Reinforcement Learning

Learning to act when the model is gone

Jinkyoo Park · KAIST

From the Bellman equation you can *solve* to the one you can only *sample*

— HANDOFF

The handoff — what Lecture 7 left us

Lecture 7 owned the world. This lecture takes it away and keeps the equation.

Where we are — the OR lineage, made data-driven

STAGES
dynamic
○ — ●

MODEL · CROSSING
model-based → data-driven
○ — ●

AGENTS
single agent
● — ○

	LINEAGE A · OR / DYNAMIC PROGRAMMING	LINEAGE B · CONTROL THEORY
MODEL-BASED (ORIGIN)	MDP & Dynamic Programming LECTURE 7	Optimal Control / Planning LECTURE 9
DATA-DRIVEN (EXTENSION)	Value-Based RL LECTURE 8	Policy-Based RL LECTURE 10

We stay in the same cell of the big map — **single agent, multiple stages** — and move along **one axis only**: from a world we are *handed* (P, R known) to a world we can only *sample*. Lecture 10 will perform the identical move on the control lineage.

What Lecture 7 gave us, stated as one object — the **Bellman optimality equation**:

$$Q^*(s, a) = \sum_{s'} P(s' \mid s, a) \left[R(s, a, s') + \gamma \max_{a'} Q^*(s', a') \right]$$

And two ways to solve it — value iteration, policy iteration. Both are exact. Both are **useless without the two highlighted terms**.

What we keep, and what we lose

The Bellman equation survives. [The model does not.](#)

Everything else follows from removing P and R . Three things break the moment they vanish:

- **The expectation.** $\sum_{s'} P(s' \mid s, a)[\dots]$ cannot be computed — we never see P , only one sampled s' at a time.
- **The improvement step.** Acting greedily on V needs P to look one step ahead. Without it, V alone cannot tell us *what to do*.
- **The table.** Even with samples, a Q -value per state–action pair does not fit in memory — nor generalize — once states are images.

Each break is one question of today. We never replace the Bellman equation; we replace the *model inside it* — first by samples, then by a function.

The translation table — our coordinate system for today

Every object of Lecture 7 has a *sampled* counterpart. Keep this in sight.

LECTURE 7 (MODEL-BASED)	LECTURE 8 (SAMPLED)	WHAT REPLACES THE MODEL
expectation $\sum_{s'} P(s' s, a)[\dots]$	one transition (s, a, r, s')	a sample average / a bootstrap
policy evaluation (solve with P, R)	MC or TD prediction	returns, or a one-step estimate
greedy on V: $\arg \max_a \sum_{s'} P[\dots]$	greedy on Q : $\arg \max_a Q(s, a)$	store Q , not V
value iteration	Q-learning	the max inside a sampled backup
full backup (all successors)	sample backup (one successor)	the law of large numbers
tabular, exact	function approximation (DQN)	a network $Q(s, a; w)$

Read column 3. Nothing here is a new *principle* — it is the same Bellman backup, with the one piece we don't own quietly swapped out. The art is in [what we swap it for](#) without the whole thing collapsing.

One honest detour — why not just learn the model?

The most literal idea: estimate P, R from data, then run Lecture 7.

$$\hat{P}(s' | s, a) = \frac{\#(s, a, s')}{\#(s, a)}, \quad \hat{R}(s, a, s') = \text{average } r \text{ over } (s, a, \cdot, s')$$

Plug in, solve the MDP. This is **model-based RL**, and it is perfectly valid.

But notice the waste: to choose actions we only ever need Q^* , yet here we first estimate a *whole transition kernel* — vastly more parameters than the thing we actually use. Why estimate P and R at all,

when we could estimate Q^* **directly**?

That question — *skip the model, estimate the value* — is the entire model-free programme. The rest of today walks it. (The model returns, deliberately, in Lecture 11.)

The roadmap — four questions

One question per Act. This strip returns at every transition — watch the highlight move.



- **Q1 — How do we evaluate a policy with no model?** Sample the whole return, or **bootstrap** from one step. (*MC vs. TD; Sutton, 1988*)
- **Q2 — How do we improve with no model?** Greedy-on- V needs P ; so learn Q , and pay for exploration with ε -greedy.
- **Q3 — Whose value are we learning?** On-policy (**SARSA**) vs. off-policy (**Q-learning**) — the cliff decides. (*Watkins, 1989*)
- **Q4 — Can it scale past a table?** Function approximation, the deadly triad, and the two tricks of **DQN**. (*Mnih et al., 2013/2015*)

— ACT 1

Evaluation without a model

Q1. Fix a policy. With no P and no R , how good is it?

The arena — from MDP to experience

Q1 Evaluate?

Q2 Improve?

Q3 Explore?

Q4 Scale?

We no longer query a model. We *act*, and a stream comes back:

$s_1, a_1, r_2, s_2, a_2, r_3, s_3, \dots, s_T$ (one episode)

- the agent chooses $a_t = \pi(s_t)$ — *decision*,
- the environment returns r_{t+1}, s_{t+1} — *the only window onto P, R* ,
- from this stream we must recover Q^* and act on it — *learning*.

Start with the humblest sub-task: **prediction**. Fix a policy π ; estimate its value V^π . No control yet — just: how good is this behavior? Everything harder is built on getting this right.

The quantity, and the two ways to estimate it

The definition has not changed since Lecture 7 — only our access to it:

$$V^\pi(s) = \mathbb{E}_\pi[G_t \mid s_t = s] = \underbrace{\mathbb{E}_\pi\left[\sum_{k \geq 0} \gamma^k r_{t+k+1} \mid s_t = s\right]}_{\text{the full return}} = \underbrace{\mathbb{E}_\pi[r_{t+1} + \gamma V^\pi(s_{t+1}) \mid s_t = s]}_{\text{one step, then bootstrap}}$$

The same value, written two ways. Each writing is a recipe for estimating it from samples:

- **Monte Carlo** reads the *left* form: wait for the episode to end, average the realized return G_t .
- **Temporal Difference** reads the *right* form: take one step, then **trust your own current estimate** of the rest.

Sample the whole thing, or sample one step and bootstrap — **the deepest fork in all of value-based RL.**

Monte Carlo prediction — average what actually happened

Visit s_t , watch the episode finish, collect the realized return $G_t = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{T-t-1} r_T$. Then nudge:

$$V(s_t) \leftarrow V(s_t) + \alpha \left[\underbrace{G_t}_{\text{target}} - V(s_t) \right]$$

- **Unbiased.** $\mathbb{E}[G_t] = V^\pi(s_t)$ exactly — the target is the real thing, not a guess.
- **No model, no bootstrap.** Each state's estimate stands alone; we never lean on neighbors.
- **The price:** must wait for the terminal state (no online updates), and the return of a long episode is a **high-variance** single draw — one unlucky tail corrupts the estimate.

MC also *wastes* the structure of the MDP: it learns each state separately and never reuses the fact that s leads to s' . That wasted structure is exactly what TD reclaims.

Temporal Difference — learn a guess from a guess

Don't wait. After a *single* transition (s_t, r_{t+1}, s_{t+1}) , update:

$$V(s_t) \leftarrow V(s_t) + \alpha \left[\underbrace{r_{t+1} + \gamma V(s_{t+1})}_{\text{TD target}} - V(s_t) \right]$$

The target contains $V(s_{t+1})$ — **our own current estimate**. We update a guess toward a slightly-better guess. This is **bootstrapping**.

- **Online, incremental, model-free.** One step is enough — ideal for long or non-terminating tasks.
- **Low variance, some bias.** The target depends on one reward plus a learned value, not a whole noisy trajectory — but it inherits whatever error $V(s_{t+1})$ currently carries.
- In practice TD usually converges *faster* than constant- α MC on stochastic tasks.

Both estimators, the same episodes



Five states, a random walk from **C**, reward 1 only on exiting right — so $V^\pi(s) = i/6$. MC and TD see the *identical* stream and differ only in the target. Watch TD's low-variance target settle while MC's realized returns keep jittering; then raise α and watch MC degrade faster.

The triangle — MC, DP, and TD

Three ways to back up a value; TD is the one that takes from both parents.

	SAMPLES (NO MODEL)	BOOTSTRAPS (USES OWN ESTIMATE)
Dynamic Programming (<i>Lec 7</i>)	no — needs full P	yes
Monte Carlo	yes	no — waits for the true return
Temporal Difference	yes	yes

TD = the model-free-ness of MC + [the bootstrapping of DP](#).

That single inheritance is why TD, not MC, becomes the engine of control. We now turn the prediction crank into a *decision-making* one.

— ACT 2

Improvement without a model

Q2. We can score a policy. How do we make it better, still without P ?

Why V is not enough — the hidden P



Lecture 7's improvement step, written out:

$$\pi'(s) = \arg \max_a \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma V(s')]$$

To turn a *value* into an *action*, you must look one step ahead — and looking ahead needs the model.


The fix is almost embarrassingly simple. Store the action-value Q instead of V :

$$\pi'(s) = \arg \max_a Q(s, a)$$

No sum, no P . Q already has the look-ahead baked in. This is *the* reason value-based RL learns Q , not V : Q is the model-free currency of decisions.

Generalized Policy Iteration, now from samples

The Lecture 7 dance is unchanged — evaluate, improve, repeat — but each half is now sampled.



One subtlety the model-based version never faced: to estimate $Q^\pi(s, a)$ from data, *every* (s, a) must actually be tried. A purely greedy policy never tries the actions it currently dislikes — and so never learns they were good.

Improvement now requires exploration. **That is the new tax.**

The tax, paid — ε -greedy, and the bandit underneath


The simplest way to keep trying everything:

$$a_t = \begin{cases} \arg \max_a Q(s_t, a) & \text{with prob. } 1 - \varepsilon \quad (\text{exploit}) \\ \text{a random action} & \text{with prob. } \varepsilon \quad (\text{explore}) \end{cases}$$

This is the **exploration–exploitation trade-off** of the n -armed bandit, now living inside every state. The bandit was RL with the dynamics deleted; here the dynamics are back, but the same dilemma — *spend on what you know, or pay to learn more* — governs every step.

- Decay $\varepsilon \rightarrow 0$ (e.g. $\varepsilon = 1/t$) and, with every (s, a) visited infinitely often, the greedy policy in the limit is optimal.
- Hold ε fixed and you keep a permanently-cautious, permanently-curious agent — sometimes exactly what a real system needs.

The tax, seen — what a greedy agent never learns



Two Q-learning agents, identical but for ϵ . The greedy one locks onto the near +1 it stumbled into first and stops looking — read its coverage figure: the actions it dislikes are never tried, so their values are never corrected. The exploring one pays a little return per episode and finds the far +5.

— ACT 3

Whose value are we learning?

Q3. The data comes from an exploring policy. Which policy's value does the update actually estimate?

Two control rules, one tiny difference



Put TD evaluation of Q together with ε -greedy improvement. Two natural updates appear, differing in *one term*:

SARSA — ON-POLICY

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

a' is the action **actually taken** next.

Q-LEARNING — OFF-POLICY

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

max over **all** next actions, taken or not.

S, A, R, S', A' — SARSA literally uses the next action it sampled. Q-learning replaces that sampled a' with the **best possible** a' . That single max is the entire on/off-policy distinction.

The 2×2 that organizes everything

Cross the two forks of the whole lecture — *bootstrap?* and *whose policy?*

	NON-BOOTSTRAP (MC)	BOOTSTRAP (TD)
On-policy	MC control	SARSA
Off-policy	off-policy MC	Q-learning

On-policy: the policy that *generates* the data is the policy we *evaluate*. SARSA learns the value of the very (exploring, ϵ -greedy) behavior it is running.

Off-policy: the behavior policy and the target policy differ. Q-learning *behaves* ϵ -greedily but *learns* the value of the pure-greedy optimal policy — it studies one policy while living as another.

Why Q-learning is "off-policy," precisely

Look at what each update assumes about the next step.

- SARSA's target uses $Q(s', a')$ where $a' \sim \pi_{\text{behavior}}$ — the **real** next action, exploration and all. It is honest about the policy it runs.
- Q-learning's target uses $\max_{a'} Q(s', a')$ — it *pretends* the next action will be greedy, regardless of the a' actually taken. The behavioral a' is collected, then **ignored**.


THE CLEAN EXTREME

act with $\epsilon = 1$, i.e. fully random

SARSA learns the value of *acting randomly*. Q-learning learns the value of the *optimal* policy — **while acting randomly**.

Because the learned target is independent of the behavior, Q converges to Q^* no matter how (sufficiently exploratory) the data was gathered.

The cliff — where the difference becomes visible




Q-learning learns the *optimal* path — along the edge — but it *acts* ϵ -greedily, and the occasional random step plunges it off the cliff: higher reward in theory, worse online. **SARSA** accounts for its own exploration and learns a **cautious** detour. Neither is "better" — the cliff teaches the design choice.

— ACT 4

Scaling past the table

Q4. Everything so far stores one number per (s, a) . What happens when the state is an image?

The wall — a table cannot hold the world



Everything so far stores one number per (s, a) . Then comes Atari:

$$\text{state} = \text{four } 84 \times 84 \text{ grayscale frames} \Rightarrow |\mathcal{S}| \approx 256^{84 \times 84 \times 4}$$

more configurations than atoms in the universe. A Q -table is impossible three times over: too many states to *visit*, too many to *store*, and — fatally — no **generalization**: a table says nothing about a state it has never seen.

The machine-learning answer: stop storing, start *approximating*.

$$Q(s, a) \approx \hat{Q}(s, a; w), \quad \text{e.g. } \hat{Q}(s, a; w) = w^\top \phi(s, a) \text{ or a neural net}$$

Now "learning" means fitting w — and similar states share answers.

Q-learning as regression — the semi-gradient step

Treat the Bellman target as a label and minimize squared error on each transition:

$$\mathcal{L}(w) = \frac{1}{2} \left(r + \underbrace{\gamma \max_{a'} \hat{Q}(s', a'; w)}_{\text{target}} - \hat{Q}(s, a; w) \right)^2 \Rightarrow w \leftarrow w + \alpha \delta \nabla_w \hat{Q}(s, a; w)$$


with TD error $\delta = r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)$.

It is called a **semi-gradient** method for an honest reason: we differentiate the prediction $\hat{Q}(s, a; w)$ but *treat the target as fixed*, even though it too depends on w . We are chasing a target that moves when we move.

In the tabular world that chase still converged. With function approximation, it can **diverge**. Tabular Q-learning's guarantees do not survive the jump — and we must understand why before trusting a network.

The deadly triad — why naive deep Q-learning blows up

Three ingredients are each harmless alone. Together they can make value estimates diverge.



Q-learning with a neural net has *all three*. The data are also strongly **correlated** (consecutive frames) and the target **shifts** on every update. Plain SGD on $\mathcal{L}(w)$ fights itself.

The breakthrough was not a new objective — it was **two engineering tricks** that tame the triad.

The triad, switched on and off

TWO STATES, ZERO REWARDS, TRUE VALUE = 0 — WHAT COULD GO WRONG

DIVERGING
 $|V(s_1)| = 5.2e+33$



Two states, every reward 0, so the true value is 0 everywhere. With all three switches on, $|V(s_1)|$ grows without bound. Turn **any one** of them off — use a table, use the return instead of a bootstrap, or update on-policy — and the same algorithm collapses to 0. (Tsitsiklis & Van Roy, 1997)

DQN — the two tricks

Mnih et al., 2013; Nature 2015

Keep Q-learning's update. Add two stabilizers.


- **Experience replay.** Store transitions (s, a, r, s') in a buffer D ; train on *random minibatches* from it. This **breaks correlation** between consecutive samples and reuses each experience many times — restoring something close to i.i.d. data.
- **Target network.** Compute the target with a *frozen* copy w^- , synced to w only every C steps:

$$\mathcal{L}(w) = \mathbb{E}_{(s,a,r,s') \sim D} \left[\left(r + \gamma \max_{a'} \hat{Q}(s', a'; w^-) - \hat{Q}(s, a; w) \right)^2 \right]$$

Now the target stops moving while we chase it — the **moving-target** pathology is suppressed.

Result: one architecture, one set of hyperparameters, **human-level play across dozens of Atari games** from raw pixels. The Bellman equation, untouched; the model, never learned; the table, replaced by a convnet.

The DQN loop — everything from today, assembled



every C steps: $w^- \leftarrow w$

Trace the lineage of each box: *act ϵ -greedy* (Act 2's exploration tax), *the max in the target* (Act 3's off-policy Q-learning), *bootstrap target* (Act 1's TD), *buffer + frozen w^-* (Act 4's triad fixes). One loop holds the whole lecture.

What comes after DQN — one slide of horizon

The two tricks opened a decade of refinements, each patching a named flaw:

- **Double DQN** — the max over-estimates; decouple action-selection from evaluation. (*van Hasselt et al., 2016*)
- **Dueling networks** — split $\hat{Q}$ into state-value + advantage; learn "which states matter" separately. (*Wang et al., 2016*)
- **Prioritized replay** — sample high-error transitions more often. (*Schaul et al., 2016*)
- **Rainbow** — combine them; the sum beats every part. (*Hessel et al., 2018*)

Every one keeps the same skeleton: *sampled Bellman backup* + Q for model-free greed + a stabilizer. None brings back the model — and none works in **continuous action spaces**, where $\max_{a'}$ itself becomes intractable.

That wall — the max over a continuum — is exactly where **Lecture 9** begins.

— CLOSING

Closing

Where we are — the cell, filled

	MODEL-BASED	MODEL-FREE
Static, single	optimization (<i>Lec 1</i>)	bandit (<i>Lec 8, intro</i>)
Dynamic, discrete, single	MDP / DP (<i>Lec 7 ✓</i>)	value-based RL (<i>Lec 8 ✓</i>)
Dynamic, continuous, single	optimal control (<i>Lec 9</i>)	policy-based RL (<i>Lec 10</i>)

One repeated move, three times over: **model** → **sample**, **wait** → **bootstrap**, **table** → **function**.

Lecture 7 could solve the Bellman equation,
because it owned the world. Lecture 8 *learned* to
solve it.

by sampling the expectation, storing Q instead of V , and replacing the table with a network it could trust.

Questions?

From a windy gridworld to human-level Atari with one update rule — the distance between *knowing* the dynamics and *never needing to*. What stays fixed through all of it is the equation Bellman wrote; we only ever changed what we were allowed to look up.

— APPENDIX

Backup slides

Complete arguments, kept out of the narrative.

Backup 1 — MC vs. TD, the bias–variance ledger

Both estimate $V^\pi(s) = \mathbb{E}_\pi[G_t \mid s_t = s]$ from samples; they differ in *what* they sample.

MONTE CARLO TARGET $G_T = \sum_K \Gamma^K R_{T+K+1}$

- *Bias*: none — $\mathbb{E}[G_t] = V^\pi(s_t)$ exactly.
- *Variance*: high — a function of **many** random rewards and transitions along the whole episode.

TD(0) TARGET $R_{T+1} + \Gamma V(S_{T+1})$

- *Bias*: present — uses the current (imperfect) estimate $V(s_{t+1})$; unbiased only at the true V^π .
- *Variance*: low — depends on **one** reward and one transition.

THE BRIDGE — N -STEP AND TD(λ)

Interpolate by bootstrapping after n steps, $G_t^{(n)} = r_{t+1} + \dots + \gamma^{n-1}r_{t+n} + \gamma^n V(s_{t+n})$; $n = 1$ is TD(0), $n = \infty$ is MC. TD(λ) averages all n geometrically. The whole family is one dial between *low-bias / high-variance* and *higher-bias / low-variance*.

Backup 2 — when does tabular Q-learning converge?

Claim (*Watkins & Dayan, 1992*): tabular Q-learning converges to Q^* with probability 1, provided

- every state–action pair (s, a) is visited **infinitely often** (the role of exploration, e.g. $\varepsilon > 0$), and
- the step sizes satisfy the Robbins–Monro conditions $\sum_t \alpha_t = \infty$, $\sum_t \alpha_t^2 < \infty$.

Why it holds (sketch). The Bellman optimality operator $(\mathcal{T}Q)(s, a) = \mathbb{E}_{s'}[r + \gamma \max_{a'} Q(s', a')]$ is a γ -contraction in the sup-norm, so it has a unique fixed point Q^* . The Q-learning update is a **stochastic approximation** of $\mathcal{T}$: each step replaces the expectation by one sample. Robbins–Monro step sizes average out the sampling noise; infinite visitation guarantees every entry keeps being corrected. Together they drive $Q \rightarrow Q^*$.

What breaks under approximation. Replace the table by $\hat{Q}(\cdot; w)$ and the update is no longer a contraction in any norm the parameters respect — the projection onto the function class can *expand* distances. That is the formal root of the deadly triad, and the reason DQN needs its two crutches.

Backup 3 — the function-approximation gradient, in full

Linear case $\hat{Q}(s, a; w) = w^\top \phi(s, a) = \sum_k w_k \phi_k(s, a)$. Per-transition squared TD error:

$$\text{Err}(w) = \frac{1}{2} \left(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w) \right)^2$$

Differentiate, holding the target fixed (semi-gradient):

$$\frac{\partial \text{Err}}{\partial w_k} = - \left(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \sum_j w_j \phi_j(s, a) \right) \phi_k(s, a)$$

$$\Rightarrow w_k \leftarrow w_k + \alpha \left(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w) \right) \phi_k(s, a)$$

In vector form $w \leftarrow w + \alpha \delta \phi(s, a)$ — the tabular update is the special case $\phi(s, a) = \mathbf{e}_{(s,a)}$ (a one-hot indicator), which is why tabular RL is "just" function approximation with the most local possible features. For a neural net, ϕ becomes $\nabla_w \hat{Q}(s, a; w)$ and the same line holds.

Backup 4 — model-based RL, the alternative we set aside

The branch we acknowledged in Act 0 and walked away from. **Estimate the model** from the same stream of experience:

$$\hat{P}(s' \mid s, a) = \frac{\#(s, a, s')}{\#(s, a)}, \quad \hat{R}(s, a, s') = \text{mean of } r \text{ over } (s, a, \cdot, s')$$

Then solve the estimated MDP with any Lecture 7 method (value / policy iteration).

- *Pro — sample efficiency*: a learned model can be "replayed" indefinitely; each real transition informs the value of *many* states through planning.
- *Con — model bias*: errors in $\hat{P}, \hat{R}$ compound through planning; you optimize for a world that is slightly wrong.
- *Con — cost*: estimating a full kernel is far more than estimating the Q^* we actually act on.

This is why model-free dominated the deep-RL era — and why **Lecture 11** returns to model-based methods only once it can make the learned model carry its weight (planning, Dyna, differentiable control).